

Final Report:

HogWildChild!

Ryan Flood and Billy Wood

15-418 Final Project

Due: May 11, 2015

Contents

1	Summary	1
2	Background	1
3	Approach	2
3.1	Back-Bone	3
3.2	Naive Approach	3
3.3	HogWild!	3
3.4	HogWildChild!	3
4	Results	4
4.1	RCV Data	4
4.2	Wikipedia Data	5
4.3	HogWildChild! and Error Rates	6
4.4	Breaking Point	7
5	Work Performed By Student	8
	References	9

1 Summary

We implemented the HogWild! algorithm, as described in a paper [1] from the University of Wisconsin-Madison, and further attempted to update the algorithm for datasets where some variables occur frequently. The algorithm is designed for performing stochastic gradient descent in parallel. We ran our code on the Intel Xeon 6-core CPUs in the GHC27-46 machines. This is the same chip used in the original HogWild! paper, and our goal was to match their results of 4x speedup.

2 Background

HogWild! is an algorithm for performing Stochastic Gradient Descent (SGD) in parallel. SGD is an optimization method for minimizing the error in a machine learning algorithm. Essentially, if you have some set of decision variables, w , you can let $E(w)$ be the error on the training set using decision variables w . Gradient descent iteratively sets $w := w - \alpha \nabla E(w)$, where α is known as the learning rate and decreases with each iteration of the algorithm (as you hone in on a sufficiently accurate solution). SGD approximates the true gradient of E using the gradient at a single data point, and sweeps through the entire training set, updating w after each example.

This is where parallelization is desired. Since each datapoint updates the decision variables, and each datapoint uses the decision variables when computing the gradient, it seems important that these changes be made serially.

The HogWild! paper gives conditions where it is okay to perform these changes in parallel. In particular, HogWild requires *sparse separable cost functions* of the form:

$$E(w) = \sum_{d \in D} e_d(w_d)$$

Where $d \in D$ denotes a small subset of the decision variables in w . Essentially, the d is a training example, which only depends on some subset of the decision variables. e_d is the error associated with that particular training sample, and so we can approximate the gradient of $E(w)$ at data point d by $\nabla e_d(w_d)$.

This provides the opportunity for a lot of data parallelism, because if data samples d_1 and d_2 use different decision variables, you can easily compute $\nabla e_{d_1}(w)$ and $\nabla e_{d_2}(w)$ in parallel and apply these updates to w .

As we sought to investigate the parallelism of this approach, we chose to focus on gradient descent in a Support Vector Machine (SVM). SVM is an algorithm for classifying samples between two different classes, represented by $\{-1, 1\}$. For a data sample z , and decision variables w , we can compute the output of the svm as $\text{sign}(w \cdot z)$. In this, the error term over all the data is:

$$\min_w \left(\sum_{d \in D} \max(1 - y_d w^T z_d, 0) \right) + \lambda \|w\|^2$$

Where $d = (z, y)$ is a data sample, with value $y \in \{-1, 1\}$. The term $\lambda \|w\|^2$ is used to prioritize small values of w , as these behave less radically on new data.

This can be changed into a sparse separable function as described by:

$$\min_w \sum_{d \in D} \left(\max(1 - y_d w^T z_d, 0) + \lambda \sum_{u \in d} \frac{w_u^2}{h_u} \right)$$

Above, $u \in d$ means the decision variables, u , used in data sample d , and h_u refers to the number of training examples which use decision variable u .

The HogWild! paper suggests that SGD can be done lock-free in parallel as long as the following three quantities are relatively small:

$$\begin{aligned} \Omega &:= \max_{d \in D} |d| \\ \Delta &:= \frac{1}{|D|} \max_{1 \leq i \leq n} |\{d \in D : i \in D\}| \\ \rho &:= \frac{1}{|D|} \max_{d \in D} |\{\hat{d} \in D : \hat{d} \cap d \neq \emptyset\}| \end{aligned}$$

Essentially, Ω says that no data sample uses too many decision variables, Δ says that no decision variable occurs in too many data samples, and ρ says that no data sample shares decision variables with too many other samples.

Unfortunately, the same thing which allows us to perform updates lock-free also prevents locality and limits SIMD execution. Namely, because the training samples are sparse, do not exhibit spatial locality with the decision variables they touch, e.g. a sample may require updating decision variable 5 and then 5000. This also seriously limits the potential of SIMD execution, since the variables being updated by a single data sample will not be adjacent in memory, thus requiring us to **gather** data from different spots in memory, which is much slower than a batch load of consecutive data.

3 Approach

We worked on HogWildChild! in three steps:

1. Naive, in which we spawn threads to do the work and each thread requests a lock before updating the decision variable vector.
2. HogWild!, in which we implement the algorithm in the HogWild! paper where each thread selects data samples at random and applies updates to the decision variables without acquiring locks.
3. HogWildChild!, which is our attempt at updating the HogWild! algorithm to allow for less sparse data samples, having some variables occur quite frequently.

3.1 Back-Bone

The basic implementation of the three iterations is the same. We worked in C++, and used PThreads for our threading.

A single master thread spawns off workers. The threads all do their work and synchronize with a barrier at the end of each *epoch*. In serial SGD, an epoch would be an iteration over the entire dataset. In our case, an epoch is time during which each thread does its assigned work of n/p samples. Upon synchronizing, the threads update α , the learning rate, such that future changes have a smaller effect on the decision variables.

The decision variables are stored as a `std::vector<double>` in memory. As computing the gradient for a particular data sample requires knowing how many times each decision variable is used, we use a `std::vector<int>` shared between the threads to store this data.

The training data is stored as an array of `sparse_examples`. A `sparse_example` has an array of `values` and another array of `indexes`. The `indexes` are the decision variables used by that example, and the `values` are the values of this sample on that decision variable. We adopted this data structure (with some modifications) from the code provided by the HogWild! research paper. The arrays in a `sparse_example` are sorted by the value in `indexes`. This is an invariant invoked by the people behind HogWild! to simplify many mathematical operations on `sparse_examples`.

We chose to work on the Intel Xeon 6-Core CPUs in GHC27-46 machines. These cores support hyperthreading, allowing for up to 12 threads to work in parallel. As such, we did all of our work on up to 12 cores, allowing all threads to run simultaneously.

3.2 Naive Approach

For our Naive approach, the thread acquires a lock before updating the decision variables.

3.3 HogWild!

The implementation of HogWild! relies on the math described in the Background section. Namely, because we know that we are dealing with sparse vectors, we update the decision variable lock-free.

3.4 HogWildChild!

For HogWildChild!, we needed a way of dealing with particularly frequent decision variables. We ultimately settled on a frequency threshold of $1/p$, where p is the number of worker threads. If a variable is below the threshold, we deemed it safe to update the decision variable in memory as soon as we saw the change. If a variable was above the threshold, we stored the pending changes in a local `UpdateVector` until enough updates had been made to the same decision variable that it was worth flushing those updates to the global copy. If a variable

occurs with frequency $\frac{k}{p}$, then we require at least k updates to that variable before flushing the updates to global memory.

An **UpdateVector** is a linked list, where each element contains an **index** and **value**. We decided to use a linked list, as it allowed us to keep the elements sorted by **index**, while still adding and removing from the list when we saw new variables with pending updates, or flushed updates to the global copy.

The $O(n)$ access time in linked lists is not an issue here. Since we only access an **UpdateVector** when we’re applying updates from a **sparse_example**, we simply iterate through the **UpdateVector** as we go through the **sparse_example**’s **indexes** array.

Finally, HogWildChild! flushes all pending updates between epochs. This prevents the **UpdateVector** from growing too big and slowing down execution.

4 Results

Our initial goal, as described in our project proposal, was to mimic the 4x speedup achieved by the HogWild! paper. We were able to do this on two different datasets.

4.1 RCV Data

The RCV dataset was provided by the people behind the original HogWild! paper. The dataset includes 800,000 training examples, and 20,000 testing examples. Data samples represent news articles, and the classifier is attempting to determine whether or not an article discusses a corporate entity.

Threads	Naive (s)	HogWild! (s)	HogWildChild! (s)
1	13.34	—	—
2	9.36	10.56	10.86
4	6.70	5.61	5.91
8	6.90	4.01	4.43
12	7.00	2.93	3.27
Train Error (12 Thread)	3.51%	3.52%	3.52%
Test Error (12 Thread)	4.27%	4.40%	4.45%

Table 1: Table of the execution times for 20 epochs of the RCV data

Figure 4.1 shows the speedup graph on the RCV data. Our 12 thread implementation easily passes the 4x goal that we set as our goal in our project proposal. It’s fairly easy to see that the Naive (lock-based) implementation reaches maximum speedup around 4 cores and then plateaus. We posit that this is because, once you reach 4 cores on this dataset, there is always a thread waiting on the lock.

We also note that our HogWildChild! implementation is slightly slower than the regular HogWild!, and that WildChild! does not improve the error rates on a 12-threaded program.




Figure 1: Graph of speedup on the RCV data

On this dataset, 0.3% of decision variables met the $1/p$ threshold for being stashed in a pending changes `UpdateVector`, with a total of 27.9% of accesses meeting this threshold. i.e, Even though only 0.3% of decision variables met the threshold, these variables occur frequently enough that they make up 27.9% of decision variables used across all training samples.

4.2 Wikipedia Data

The Wikipedia dataset consists of approximately 120,000 training samples and 10,000 testing samples. We made this dataset by downloading Wikipedia’s database in German and Portuguese, and converting articles into the 3-grams therein. As shown in Table 4.2, just by looking at the 3-grams we were able to determine language with approximately 98% accuracy.

As in the RCV dataset, we were easily able to achieve the targeted 4x speedup over the serial approach.

In this dataset, we see that the Naive approach does not plateau nearly as dramatically as in the RCV data. This is likely because samples in our Wiki dataset were smaller than in the RCV dataset, and so required holding the lock for less time.

Further, the HogWildChild! implementation is still slower than HogWild!, and no more accurate. On this dataset, 0.23% of decision variables pass the $1/p$ threshold on 12 threads. However, 39.81% of decision variables used across all training samples pass by this threshold. This means that 40% of the time, when

Threads	Naive (s)	HogWild! (s)	HogWildChild! (s)
1	9.51	—	—
2	5.31	5.56	5.72
4	3.31	2.94	3.13
8	2.57	2.20	2.50
12	2.22	1.63	1.88
Train Error (12 thread)	1.96%	1.96%	1.98%
Test Error (12 thread)	2.21%	2.18%	2.24%

Table 2: Table of the execution times for 20 epochs of the Wikipedia data




Figure 2: Graph of speedup on the Wikipedia language recognizer

updating a decision variable it’s likely that another thread is currently using the same variable.

4.3 HogWildChild! and Error Rates

Our goal with HogWildChild! was to come up with an approach which better dealt with frequent variables than the HogWild algorithm. However, we saw that the HogWild! algorithm achieved the same error on both of our datasets as the serial algorithm. As such, we were unable to come up with a parallel algorithm that could actually achieve better accuracy than HogWild!, and we were surprised by how well HogWild! performed even when 40% of the variables used occurred over the $1/p$ threshold.

4.4 Breaking Point

HogWildChild! uses a $1/p$ threshold to determine which decision variables should have changes stashed until sufficiently many changes had been seen. As such, increasing the number of threads will increase the percentage of decision variables which are over the threshold. The unix.andrew.cmu.edu machines have 20 cores, with up to 40 processes hyperthreaded. When run on the RCV data, we see:




Figure 3: Graph of speedup on the Wikipedia language recognizer

In the 40-threaded run on the RCV data, the 1.4% of decision variables above the $1/p$ threshold accounted for nearly 60% the decision variables used. Further, since the threshold is so much smaller, decision variables need to stay in the `UpdateVector` until they’ve witnessed more changes. This accounts for the serious slow-down of HogWildChild! around 20 threads.

Figure 4.4 shows a similar result, with HogWildChild! slowing down after 20 threads. At 40 threads, 64% of decision variables used are above the $1/p$ frequency threshold.

Interestingly, Figure 4.4 shows HogWild! plateauing around 12 threads. This is likely due to the memory bandwidth constraints, as at this point HogWild! is reading and writing faster than unix.andrew.cmu.edu’s memory bus can deal with.




Figure 4: Graph of speedup on the Wikipedia language recognizer

5 Work Performed By Student

Equal work was performed by both project members.

References

- [1] Feng Niu, Benjamin Recht, Christopher Ré, and Stephen J. Wright. *Hog-Wild!: A Lock-Free Approach to Parallelizing Stochastic Gradient Descent*, Computer Sciences Department, University of Wisconsin-Madison, 2011.
<http://www.eecs.berkeley.edu/~brecht/papers/hogwildTR.pdf>